



RV College of
Engineering®

Department of Computer Science & Engineering

Distributed
Computing



Parallel
Computing

OpenMP, MPI

Dr. Minal Moharir

Go, change the world®



RV College of
Engineering®

What is OpenMP?

De-facto standard API for writing shared memory parallel applications in C, C++, and Fortran

Consists of:

1. Compiler Directives
2. Run time Libraries
3. Environment Variables

Specifications are mentationed by OpenMp Architecture Review Board(<http://www.openmp.org>)

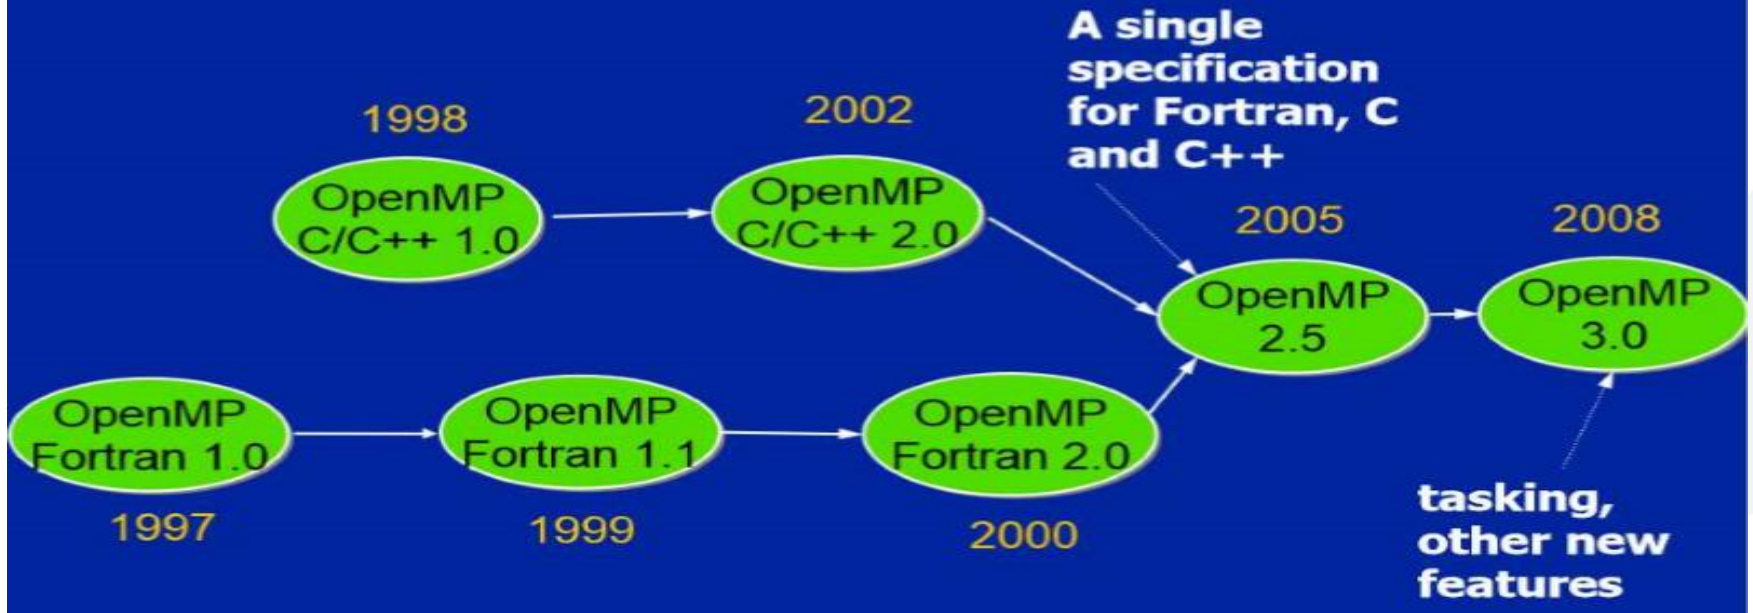
Version 4.5 has been released by in Nov. 2015

Go, change the world®



OpenMP?

OpenMP Release History



Go, change the world®



When Compiler cannot find parallelism

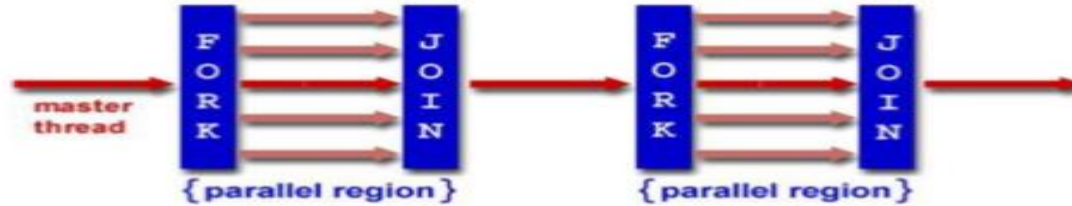
Use explicit PARALLELIZATION-OpenMp

Go, change the world®



Memory Model

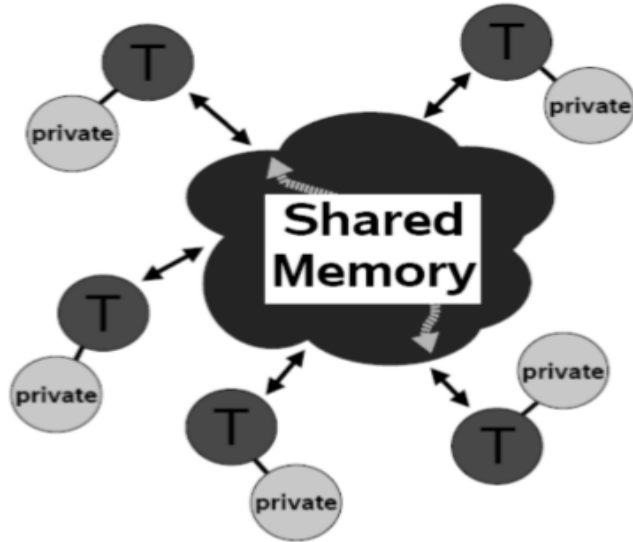
- **Shared Memory, Thread Based Parallelism**
- **Explicit Parallelism**
- **Fork - Join Model**



- **Compiler Directive Based**
- **Nested Parallelism Support**
- **Dynamic Thread**
- **Memory Model : Flush often**

Go, change the world®

Memory Model



- Data is **private** or **shared**.
- All threads have **access** to **same globally shared memory**.
- Shared data accessible by all threads.
- Private accessed only **by owned threads**.
- Data transfer is transparent to programmer.
- **Synchronization** takes place, but it is almost **implicit**.

Go, change the world®



Compilation

GNU	gcc	-fopenmp
Linux Opteron/Xeon	g++	
IBM Blue Gene	g77	
	gfortran	

GNU Compiler Example :

- **gcc -o omp_helloc -fopenmp omp_hello.c**

Go, change the world®



RV College of
Engineering®

Advantages of OpenMP

- **Good performance and scalability**
 - ✓ If you do it right
- **De-facto and mature standard**
- **An OpenMP program is portable**
 - ✓ Supported by a large number of compilers
- **Requires little programming effort**

Go, change the world®



Components of OpenMP

Compiler Directives

- **Parallel Construct**
- **Work Sharing**
- **Synchronization**
- **Data Environment**
 - ✓ private
 - ✓ first private
 - ✓ last private
 - ✓ shared
 - ✓ reduction

Environment Variables

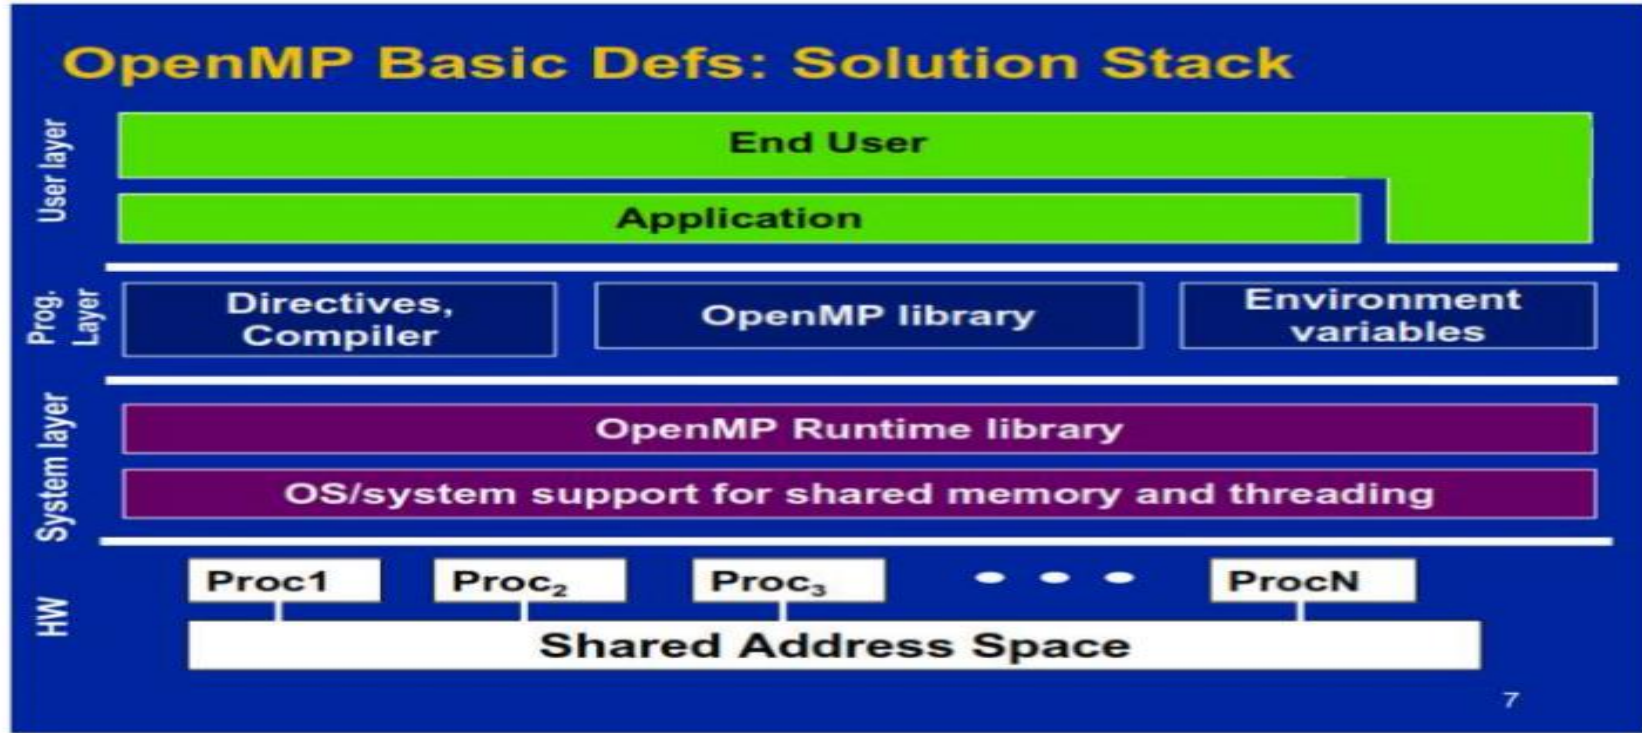
- **Number of threads**
- **Scheduling Type**
- **Nested parallelism**
- **Dynamic Thread**
- **Adjustment**

Runtime Library routines

- **Number of threads**
- **Thread ID**
- **Dynamic thread**
- **adjustment**
- **Nested parallelism**



OpenMP Stack



Go, change the world®



OpenMp Directives

❖ **#pragma omp directive-name [clause, clause..] new-line**

- **Eg: #pragma omp parallel default(shared) private(beta,pi)**

❖ **General Rules:**

- **Case sensitive**
- **Compiler Directives follow C/C++ standards**
- **Only one directive-name to be specified per directive**
- **Each directive applies to at most one succeeding statement.**
- **Use ("\\") for continuing on succeeding lines.**

Go, change the world®



Restrictions

- A parallel region must be a **structured block** that does not span multiple routines or code files
- It is **illegal to branch into or out** of a parallel region
- Only a **single IF clause** is permitted
- Only a **single NUM_THREADS** clause is permitted



How many Threads?

- **The number of threads in a parallel region is determined by the following factors, in order of precedence:**
 - ✓ **Evaluation of the IF clause**
 - ✓ **Setting of the NUM_THREADS clause**
 - ✓ **Use of the omp_set_num_threads() library function**
 - ✓ **Setting of the OMP_NUM_THREADS environment variable**
 - ✓ **Implementation default - usually the number of CPUs on a node, though it could be dynamic (see next bullet).**
- **Threads are numbered from 0 (master thread) to N-1**



Data Scoping

The OpenMP Data Scope Attribute Clauses are used to explicitly define how variables should be scoped. They include:

- ✓ PRIVATE
- ✓ FIRSTPRIVATE
- ✓ LASTPRIVATE
- ✓ SHARED
- ✓ DEFAULT
- ✓ REDUCTION

Data Scope Attribute Clauses are used in conjunction with several directives (PARALLEL, DO/for, and SECTIONS) to control the scoping of enclosed variables



Data Scoping

❖ private clause

- This declares variables in its list to be private to each thread
- Format

✓ private (list)

Eg: int B = 10;

 #pragma omp parallel private(B)

 B = ... ;

- A private un-initialised copy of B is created before the parallel region begins
- B value is not the same within the parallel region as outside

Go, change the world®



Data Scoping

❖ firstprivate clause

- **Format**

- ✓ first private (list)

- ✓ Eg: `int B;`

- `B = 10;`

- `#pragma omp parallel firstprivate(B)`

- `B = B + ... ;`

- A private **initialized copy** of B is created before the parallel region begins
- The copy of each thread **gets the same value**



Data Scoping

❖ SHARED Clause

- ✓ A shared variable **exists in only one memory location** and all threads can read or write to that address

- **Format**

- ✓ shared (list)

❖ DEFAULT Clause

- ✓ Specify default scope for all variables in the lexical extent.

- **Format**

- ✓ default (shared | none)

❖ LASTPRIVATE Clause

- ✓ Value from the **last** loop iteration **assigned the original variable** object.

Go, change the world®



Reduction Clause

✓ **Variables which needed to be shared & modified by all the processors**

- **Format**

- ✓ reduction (operator: list)

- **Example**

```
total = 0.0;
# pragma omp parallel for private ( i, p ) /
  shared ( n, x ) reduction ( +: total )

for ( i = 0; i < n; i++ )
{
    p = ( ( x[i] - 7 ) * x[i] + 4 ) * x[i] - 83;
    total = total + p;
}
```

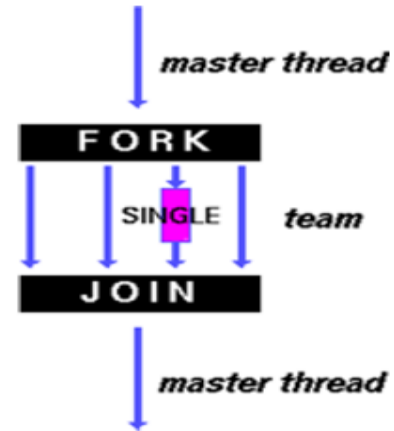
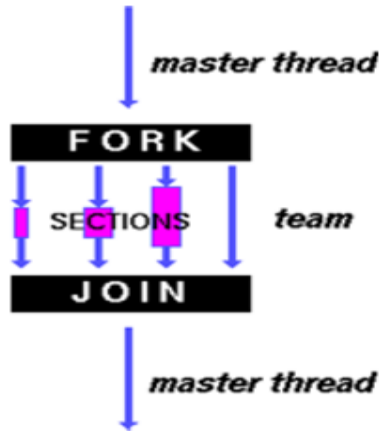
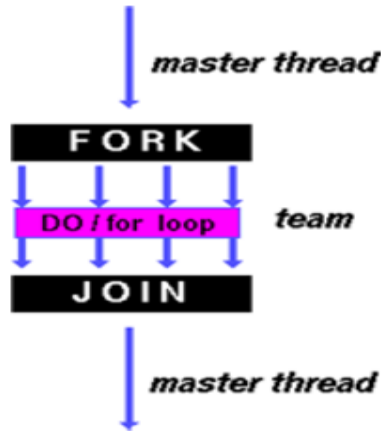
Symbol	Meaning
+	Summation
-	Subtraction
*	Product
&	Bitwise AND
	Bitwise OR
^	shift
&&	Logical AND
	Logical OR

Go, change the world®



Reduction Clause

- **FOR** - data parallelism
- **SECTIONS** - functional parallelism
- **SINGLE** - serializes a section of code



Go, change the world®



RV College of
Engineering®

Program Programming Paradigm

```
#include <stdio.h>
#include <mpi.h>
main(int argc, char
**argv)
{   int ierr;
    ierr = MPI_Init(&argc,
&argv);
    printf("Hello world\n");
    ierr =
MPI_Finalize();
}
```

```
mpicc hello.c
-o hello
mpirun -np
4 hello
```

Go, change the world®



RV College of
Engineering®

OpenMP

```
#include <omp.h>
#include <stdio.h>
int main (int argc, char
*argv[])
{
#pragma omp parallel
    printf("Hello World");
}
```

Compile:
gcc –
fopenmp
hello.c

Go, change the world®



OpenMP

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
```

```
int main (int argc, char *argv[])
{
    int nthreads, tid;
```

```
/* Fork a team of threads giving them their own copies of
variables */
```

```
#pragma omp parallel private(nthreads, tid)
{
```

```
/* Obtain thread number */
tid = omp_get_thread_num();
printf("Hello World from thread = %d\n", tid);
```

```
/* Only master thread does this */
if (tid == 0)
```

```
{
    nthreads = omp_get_num_threads();
    printf("Number of threads = %d\n", nthreads);
}
```

Go, change the world®

```
} /* All threads join master thread and disband */
```

Compile: gcc – fopenmp hello.c



MPI: Message Passing Interface

- **A message passing library specification**
- **Message passing among processes in parallel computing**
- **Meant for clusters and network of workstations**
- **Message Passing Standard for Parallel Programs**
 - MPI Implementation left to individual vendors.
 - Most commonly supports C, Fortran, C++ programs



MPI: Message Passing Interface

OpenMPI Implementation <https://www.open-mpi.org/>

- ▶ Open source (Git)
- ▶ Collaborators (partial)
 - ▶ Auburn, Wisconsin at La Crosse, Michigan
 - ▶ Los Alamos National Lab, Oak Ridge National Lab, Sandia National Lab
 - ▶ Amazon, AMD, ARM, Broadcom, Cisco, Facebook, Fujitsu, IBM, Intel, nVIDIA, Oracle

Go, change the world®



MPI: Message Passing Interface

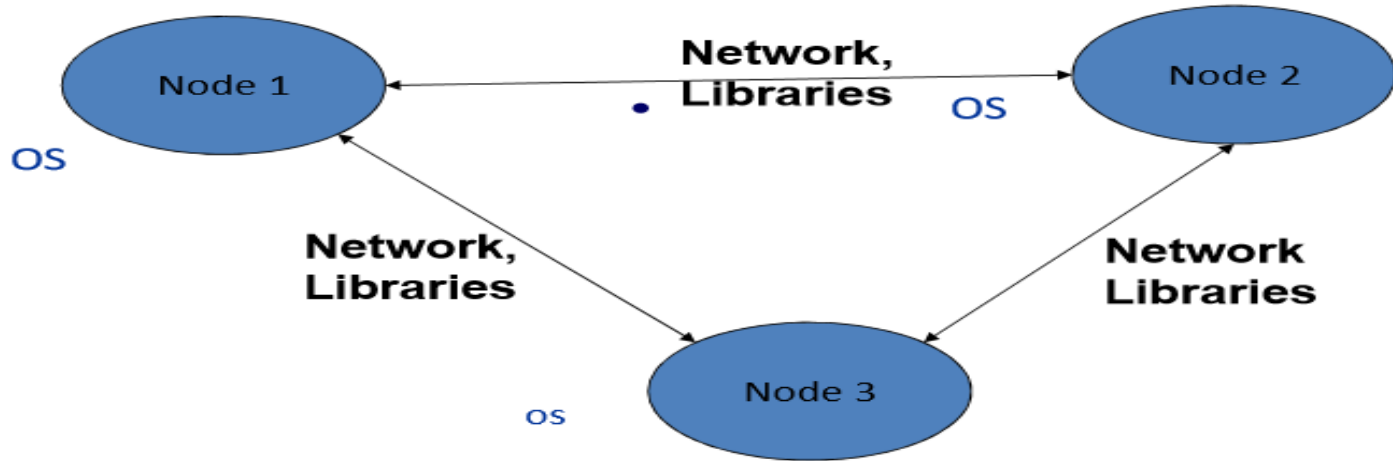
MPICH Implementation <http://www.mpich.org/>

- ▶ Open source (Git)
- ▶ 'CH' from *Chameleon* portability system
- ▶ MPICH1 (1992), MPICH2 (2001), MPICH3 (2012)
- ▶ Collaborators (partial)
 - ▶ University of British Columbia, Ohio State
 - ▶ Microsoft, IBM, Cray

Go, change the world®



MPI: Message Passing Interface



- Multiple Nodes connected together to form a Computer Cluster

Go, change the world®



RV College of
Engineering®

MPI: Message Passing Interface

MPI include file

...

Initialize MPI environment

...

Do work and make message passing calls

...

Terminate MPI Environment

Go, change the world®



RV College of
Engineering®

OpenMP

```
#include <omp.h>
#include <stdio.h>
int main (int argc, char
*argv[])
{
#pragma omp parallel
num_threads(8)
printf("Hello World");
}
```

Go, change the world®

Compile:
gcc –
fopenmp
hello.c



RV College of
Engineering®

Program Programming Paradigm

```
#include <stdio.h>
#include <mpi.h>
main(int argc, char **argv)
{
    int ierr, num_procs, my_id;
    ierr = MPI_Init(&argc, &argv);
    /* find out MY process ID, and how
    many processes were started. */
    ierr =
MPI_Comm_rank(MPI_COMM_WORLD, &my_id);
    ierr =
MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
}
```

```
printf("Hello world
I'm process %i out of
processes\n",
    my_id, num_procs);
ierr = MPI_Finalize;
```

Go, change the world®



RV College of
Engineering®

Program Programming Paradigm

```
#include <stdio.h>
#include <mpi.h>
main(int argc, char **argv)
{
    int ierr, num_procs, my_id;
    ierr = MPI_Init(&argc, &argv);
    /* find out MY process ID, and how
    many processes were started. */
    ierr =
MPI_Comm_rank(MPI_COMM_WORLD, &my_id);
    ierr =
MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
}
```

```
printf("Hello world
I'm process %i out of
processes\n",
    my_id, num_procs);
ierr = MPI_Finalize;
```

Go, change the world®



RV College of
Engineering®



THANK YOU

Go, change the world®