



PARALLEL ARCHITECTURE & DISTRIBUTED PROGRAMMING(21CS71)

DR. MINAL MOHARIR



Introduction to Parallel Computing

Objective : Motivation to learn Parallel Computing



Serial Computing




If **1 Man** takes **10 hours** to complete the job; then
How long will it take if **10 Men** work **together**

Parallel Computing

**THINK
PARALLEL**

PARALLEL Computing!

**सी डैक
CDAC**

A black and white illustration showing a grid of 21 stylized human figures, arranged in 3 rows and 7 columns. Each figure is holding a shovel and is in the process of digging a small pile of dirt. This visual metaphor represents parallel processing, where many tasks are performed simultaneously.

- **Save wall clock time**
- **Solve larger problems**
- **Provide concurrency**
 - (do multiple things at the same time)

Parallelism is the future of computing!


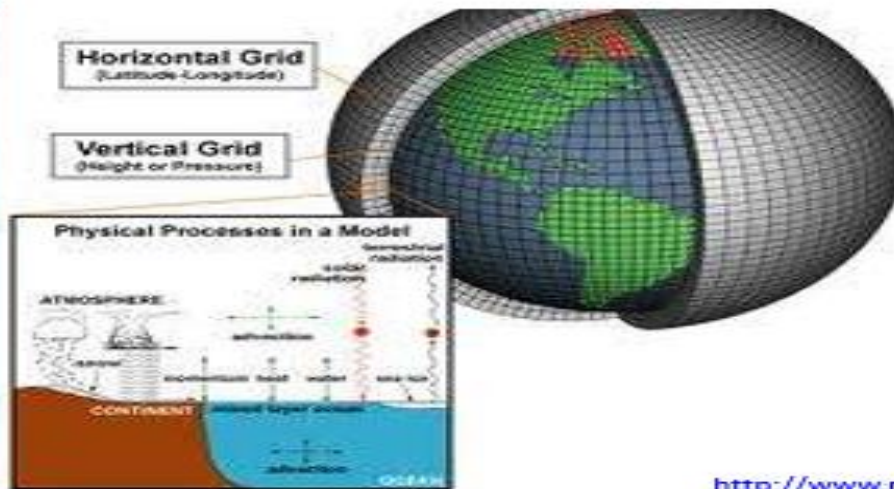
Applications

Application Drivers for HPC

सी डैक
CDAC

Weather Forecasting

- Atmosphere modeled by dividing it into 3-dimensional cells.
- Calculations of each cell repeated many times to model passage of time.



http://www.nsf.gov/news/special_reports/clouds/images/nsf_clouds_computing.pdf

Applications

THINK
PARALLEL

Weather Forecasting Example

सी डैक
CDAC

- Consider atmosphere to be divided into cells of size 1 mile \times 1 mile \times 1 mile to a height of 10 miles
- Global region can be schemed as 2×10^9 cells.
- Each cell has many physical variables (temperature, pressure, humidity, wind speed and direction, etc) to be computed
- Suppose each calculation requires 200 floating point operations. In *one time step*, 10^{11} floating point operations necessary.
- To forecast the weather over 7 days using 1-minute intervals, takes $(7 * 24 * 60 * 10^{11}) = 10080 * 10^{11} = 10^{15}$ FP operations
- a computer operating at 1Gflops (10^9 floating point operations/s) takes 10^6 seconds or over 10 days.
- To perform calculation in 5 minutes requires computer operating at 3.4 Tflops (3.4×10^{12} floating point operations/sec).



UNITS

THINK
PARALLEL

Units of Measurement in HPC

सी डैक
CDAC

- **Mflop/s** 10^6 flop/sec
- **Gflop/s** 10^9 flop/sec
- **Tflop/s** 10^{12} flop/sec
- **Pflop/s** 10^{15} flop/sec
- **Mbyte** 10^6 byte (also $2^{20} = 1048576$)
- **Gbyte** 10^9 byte (also $2^{30} = 1073741824$)
- **Tbyte** 10^{12} byte
(also $2^{40} = 10995211627776$)
- **Pbyte** 10^{15} byte
(also $2^{50} = 1125899906842624$)


Applications

THINK
PARALLEL

Applications of Parallel Computing

सी डैक
CDAC

- Prediction of weather, climate, global changes
- Modeling motion of astronomical bodies
- Human genome
- Challenges in materials sciences
- Semiconductor design
- Superconductivity
- Structural biology
- Design of drugs
- Challenges in transportation
- Vehicle dynamics
- Nuclear fusion
- Enhanced oil and gas recovery
- Computational ocean sciences
- Speech
- Vision
- Visualization and graphics



Where are we Today?

Cost, Size, Performance

Computer Technology Today have made incredible progress in roughly 65-years after first general purpose electronic computer was created.



Where are we Today?

THINK
PARALLEL

Computer Technology Today

सी डैक
CDAC

Today's Desk Top Computer for US\$ 500 more powerful than a

A modern desktop computer setup including a black tower unit, a large monitor displaying a blue abstract image, a keyboard, and a mouse.


Powerful : Performance, Main Memory & Disk Storage

High Performance Computing (HPC) System two decades ago for US\$ 1 Million

A large room filled with numerous black, floor-standing computer cabinets arranged in rows, representing a High Performance Computing (HPC) system from two decades ago.

Where are we Today?

Issue: increase in clock speed, increase Power, Heat



Definition



High Performance Computing



- **Definition 1 (Wikipedia)**
 - High Performance Computing (HPC) uses Supercomputers and Computer Clusters to solve advanced computational problems. Today, Computer Systems approaching the Teraflops- region are counted as HPC-Computers
- **Definition 2**
 - High Performance Computing (HPC) is the use of Parallel Processing for running advanced application programs efficiently, reliably and quickly. The term applies especially to systems that function above a TeraFlop or 10^{12} Floating-Point Operations Per Second.

Benchmarking

THINK
PARALLEL

TOP500

सी डैक
CDAC

Jack Dongarra, H. Simon, E. Strohmaier and H. Meuer

- Listing of the 500 most powerful Supercomputers in the World
- Based on **LINPACK Benchmark** : Measure of System's **Floating Point Computing Power**
- How **fast** a Computer solves a Dense N by N System of Linear Equations

$$Ax=b, \text{ dense problem}$$



- Result is reported as **Rate of Execution in MFlops /Sec**
 - Updated **twice a year**
- Supercomputing Conference (SC) in United States in November**
International Supercomputing Conference (ISC) in Europe in June



www.top500.org June 2024

Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
1	Frontier - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE DOE/SC/Oak Ridge National Laboratory United States	8,699,904	1,206.00	1,714.81	22,786
2	Aurora - HPE Cray EX - Intel Exascale Compute Blade, Xeon CPU Max 9470 52C 2.4GHz, Intel Data Center GPU Max, Slingshot-11, Intel DOE/SC/Argonne National Laboratory United States	9,264,128	1,012.00	1,980.01	38,698
3	Eagle - Microsoft NDv5, Xeon Platinum 8480C 48C 2GHz, NVIDIA H100, NVIDIA Infiniband NDR, Microsoft Azure United States	2,073,600	561.20	846.84	
4	Supercomputer Fugaku - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,630,848	442.01	537.21	29,899
5	LUMI - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE EuroHPC/CSC Finland	2,752,704	379.70	531.51	7,107
6	Alps - HPE Cray EX254n, NVIDIA Grace 72C 3.1GHz, NVIDIA GH200 Superchip, Slingshot-11, HPE Swiss National Supercomputing Centre (CSCS) Switzerland	1,305,600	270.00	353.75	5,194
7	Leonardo - BullSequana XH2000, Xeon Platinum 8358 32C 2.6GHz, NVIDIA A100 SXM4 64 GB, Quad-rail NVIDIA HDR100 Infiniband, EVIDEN EuroHPC/CINECA Italy	1,824,768	241.20	306.31	7,494

Principles of CA Design



Take Advantage of Parallelism

- e.g. multiple processors, disks, memory banks, pipelining, multiple functional units

Principle of Locality

- Reuse of data and instructions
- Temporal locality states that recently accessed items are likely to be accessed in the near future.
- Spatial locality says that items whose addresses are near one another tend to be referenced close together in time.

Focus on the Common Case

- The instruction fetch and decode unit of a processor may be used much more frequently than a multiplier, so optimize it first
- Addition of 2-nos., Overflow



Dependency Analysis

- **Their order of execution must not matter!**

- **Example 1**

a=1;

b=2;

- **Example 2**


a=2;

b=a;

Amdahl's Law

Inherently serial: I/O

Speedup





Program Programming Paradigm

OpenMp

MPI

CUDA

OpenACC

OpenCL



Program Programming Paradigm


```
7 #include <stdio.h>
6
5 void hello_cuda() {
4     printf("Hello Cuda\n");
3 }
2
1 int main(int argc, char **argv) {
    hello_cuda();
1     return 0;
2 }
```



Program Programming Paradigm

```
8 #include <stdio.h>
7
6 __global__ void hello_cuda() {
5     printf("Hello Cuda\n");
4 }
3
2 int main(int argc, char **argv) {
1     hello_cuda<<<1,1>>>>();
9     cudaDeviceSynchronize();
1
2     return 0;
3 }
```

Pro



F

2x2

	B 0 0	B 0 1
	B 1 0	B .

Program Programming Paradigm

2x2

T00	01	00	01
10	11	10	11



Program Programming Paradigm

```
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ mv hello_cuda.c
mv: missing destination file operand after 'hello_cuda.c'
Try 'mv --help' for more information.
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ mv hello_cuda.cu hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ls
hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ gcc hello_cuda.c -o hello_cuda
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda
Hello Cuda
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ mv hello_cuda.c hello_cuda.cu
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.cu
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.cu -o hello_cuda
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.cu
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.cu -o hello_cuda
```



Program Programming Paradigm

```
#include <stdio.h>

__global__ void hello_cuda() {
    printf("Hello Cuda\n");
}

int main(int argc, char **argv) {
    hello_cuda<<<2,2>>>();
    cudaDeviceSynchronize();

    return 0;
}
```



```
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ mv hello_cuda.c hello.c
nv: missing destination file operand after 'hello_cuda.c'
Try 'mv --help' for more information.
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ mv hello_cuda.c hello.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ls
hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ gcc hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda
Hello Cuda
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ mv hello_cuda.c hello.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda
Hello Cuda
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda
Hello Cuda
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.c
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda
[[[Aneuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda
Hello Cuda
Hello Cuda
Hello Cuda
Hello Cuda
Hello Cuda
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$
```



Program Programming Paradigm

```
5 #include <stdio.h>
4
3 __global__ void hello_cuda() {
2     printf("Hello Cuda\n");
1     printf("Block Index X: %d, Block Index Y: %d, Thread Index X: %d, Thread Index Y: %d\n",
            blockIdx.x, blockIdx.y, threadIdx.x, threadIdx.y);
1 }
2
3 int main(int argc, char **argv) {
4     hello_cuda<<<2,2>>>>();
5     cudaDeviceSynchronize();
6
7     return 0;
8 }
```




Program Programming Paradigm

```
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda  
Hello Cuda  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda  
Hello Cuda  
Hello Cuda  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.  
^[[Aneuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda  
Hello Cuda  
Hello Cuda  
Hello Cuda  
Hello Cuda  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nv hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ nvcc hello_cuda.  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$ ./hello_cuda  
Hello Cuda  
Hello Cuda  
Hello Cuda  
Hello Cuda  
Block Index X: 1, Block Index Y: 0, Thread Index X: 0, Thread Index Y: 0  
Block Index X: 1, Block Index Y: 0, Thread Index X: 1, Thread Index Y: 0  
Block Index X: 0, Block Index Y: 0, Thread Index X: 0, Thread Index Y: 0  
Block Index X: 0, Block Index Y: 0, Thread Index X: 1, Thread Index Y: 0  
neuralnine@pop-os:~/Documents/Programming/NeuralNine/Python/Current$
```




THANK YOU